


Distributed Systems

Definition of a Distributed System

- A distributed system is a collection of independent computers that appear to the users of the system as a single computer
- Two aspects:
 - Hardware: autonomous machines
 - Software: the users think of the system as a single computer

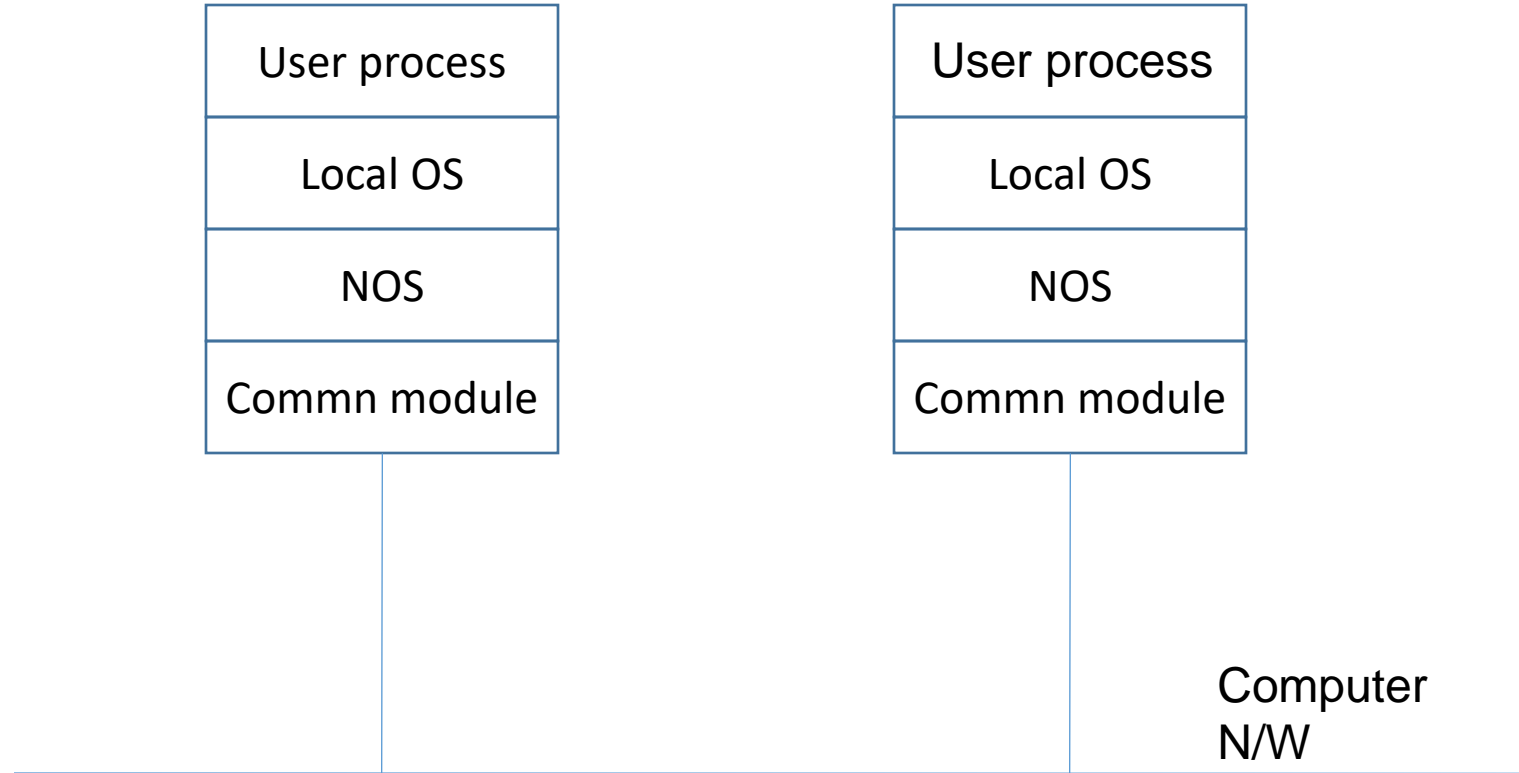
A Distributed System



Motivation

- Resource sharing
 - sharing and printing files at remote sites
 - processing information in a distributed database
 - using remote specialized hardware devices
- Computation speedup – *load sharing*
- Reliability – detect and recover from site failure, function transfer, reintegrate failed site
- Communication – message passing

Network OS



Network OS

- Users are aware of multiplicity of machines. Access to resources of various machines is done explicitly by:
 - Remote logging into the appropriate remote machine.
 - Transferring data from remote machines to local machines, via the File Transfer Protocol (FTP) mechanism.

Characteristics

- Local OS
- Own computer or designated computer or remote login
- Explicit file transfer commands
- Little or no fault tolerance

Communication and information sharing

Shared Global File system

Use of file servers

Hierarchical file system

Distributed OS

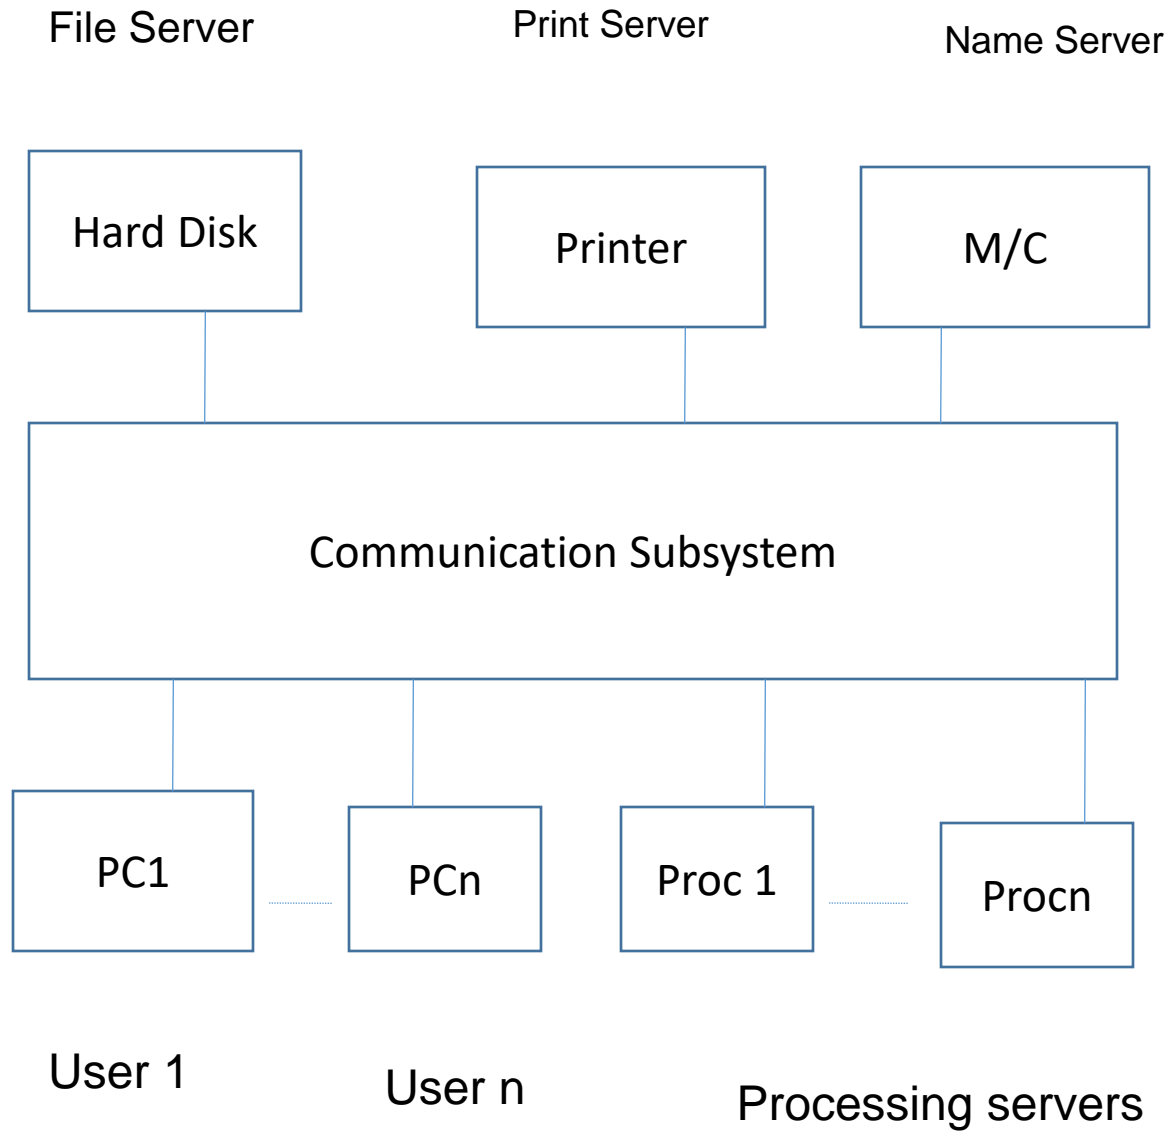
Users not aware of multiplicity of machines. Access to remote resources similar to access to local resources.

- Data Migration – transfer data by transferring entire file, or transferring only those portions of the file necessary for the immediate task.
- Computation Migration – transfer the computation, rather than the data, across the system.

- Process Migration – execute an entire process, or parts of it, at different sites.
 - Load balancing – distribute processes across network to even the workload.
 - Computation speedup – subprocesses can run concurrently on different sites.
 - Hardware preference – process execution may require specialized processor.
 - Software preference – required software may be available at only a particular site.
 - Data access – run process remotely, rather than transfer all data locally.

A DOS should

- Control n/w resource allocation
- Provide virtual computer
- Hide the distribution of resources
- Provide protection mechanisms
- Provide secure communication



Characteristics

- Global system wide OS
- Dynamic allocation of processes to CPU's
- File placement managed by OS
- Fault tolerance
- Single Global IPC mechanism

Middleware based Distributed Systems

- Middleware- Additional layer on top of NOS implementing general purpose services and it hides more or less the heterogeneity of the collection of underlying platforms.

Middleware Services:

Communication facilities

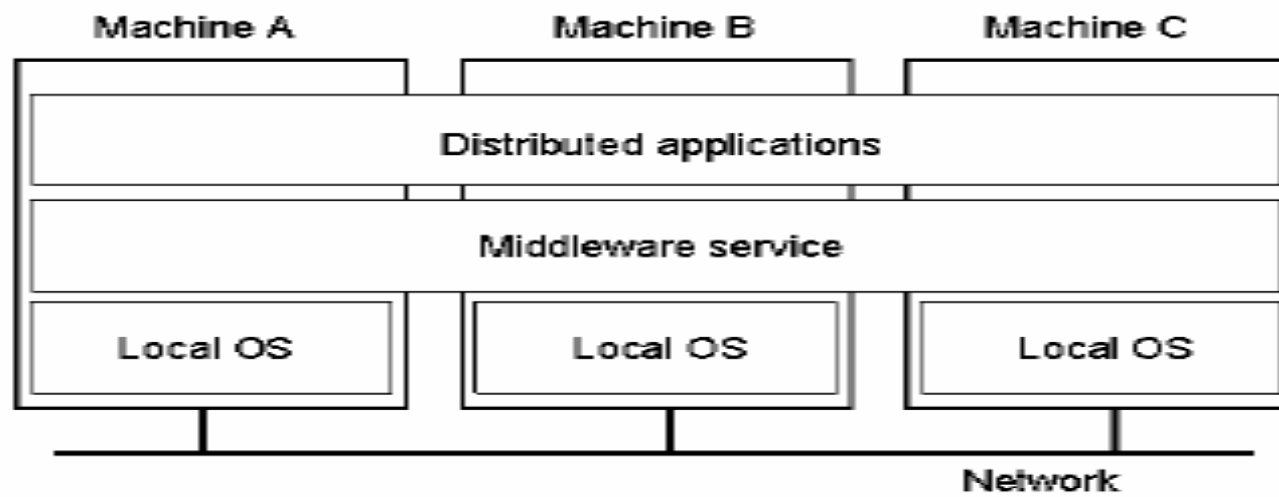
Naming

Persistence

Security

Openness

Scalability



A distributed system organized as middleware.

Goals

- A distributed system should easily connect users to resources
- Why sharing?
 - To make it easy for users to access remote resources, and to share them with other users in a controlled way
- Reason for sharing resources: economics
 - Typical resources : Printers, computers, storage facilities, data, files
 - Easier to collaborate and exchange information
 - Internet, groupware
- Problems with sharing
 - Security is becoming more and more important as connectivity and sharing increase
 - Tracking communication to build up a preference profile of a specific use
 - Unwanted Communication

Goals

- It should hide the fact that resources are distributed across a network
- It should be open
- It should be scalable

Design Issues

- Transparency – concealment from the user and the application programmer of the separation of components in a distributed systems so that the system is perceived as a whole rather than as a collection of independent components
- Flexibility- ease of modification and ease of enhancement
- Openness – The openness of a computer system is the characteristics that determines whether the system can be extended and re-implemented
- Scalability – A system is scalable, if it will remain effective when there is a significant increase in the numbers of resources and number of users
- Heterogeneity – Variety and differences
- Security- CIA
 - Confidentiality – protection against disclosure to unauthorized individuals
 - Integrity – Protection against alteration or corruption
 - Availability – Protection against interference with the means to access the resources.

Transparency in a Distributed System

Access transparency

- Hide differences in data representation and how a resource is accessed
- Intel (little endian format- *least* significant byte in the smallest address)/Sun SPARC (big endian- most significant byte in the smallest address) (order of bytes)
- OS with different file name conversions

Transparency in a Distributed System

- **Location transparency**

Hide where a resource is located importance of naming, e.g., URLs

- **Migration transparency**

Hide that a resource may move to another location

- **Relocation transparency**

Hide that a resource may move to another location while in use
example, mobile users

Transparency in a Distributed System

- **Replication transparency**

Hide that a resource is replicated subsumes that all replicas have the same name (and thus location transparency)

- **Concurrency transparency**

Hide that a resource may be shared by several competitive users leave the resource in a consistent state more refined mechanism transactions

Transparency in a Distributed System

- **Failure transparency**

Hide the failure and recovery of a resource

- **Persistent transparency**

Hide whether a (software) resource is in memory or Disk

Important problem: inability to distinguish between a dead resource and a painfully slow one

Design Issues: Transparency

Transparency	Description
Access	Hide differences in data representation and how a resource is accessed
Location	Hide where a resource is located
Migration	Hide that a resource may move to another location
Relocation	Hide that a resource may be moved to another location while in use
Replication	Hide that a resource is replicated
Concurrency	Hide that a resource may be shared by several competitive users
Failure	Hide the failure and recovery of a resource

Openness

Open distributed system

- Be able to interact with services from other open systems, irrespectively of the underlying environment
- Offers services according to standard rules that describe the **syntax and the semantics of these services**

Rules formalized in protocols

- Services specified through interfaces (described in an Interface Definition Language (IDL) (but only the syntax part)
- Neutral and complete specifications (with regards to a potential implementation)

Openness

- Interoperability: to what extent can work together
- Portability: to what extent an application developed for A can be executed on B that implements the same interface with A

Openness

- Separate Policy from Mechanism
 - A system organized as a collection of relatively small and easily replaceable or adaptable components
 - Provide definitions of the internal parts of the system as well
 - A distributed system provides only **mechanisms**.

Policies are specified by applications and users.

Example policies:

- What level of consistency do we require for client-cached data?
- Which operations do we allow downloaded code to perform?
- Which QoS requirements do we adjust in the face of varying bandwidth?
- What level of secrecy do we require for communication?

Scalability

- size (number of users and/or processes)
- geographical (maximum distance between nodes)
- administrative (number of administrative domains)
- Solution: Powerful server

Scalability Problems

Centralized algorithms - Doing routing based on complete information

Centralized data - A single on-line telephone book

Centralized services - A single server

Major problem : Single point of failure.

Decentralized algorithms

- No complete information about the system state
- Make decision only on local information
- Failure of one machine does not ruin the algorithm
- No assumption of a global clock

Scalability

- Synchronous communication
 - In WAN, Unreliable and point-to-point
- Geographical scalability
 - How to scale a distributed system across multiple, independent administrative domains:
conflicting policies with respect to resource usage (and payment), management and security
- Expand to a new domain
 - Protect itself against malicious attacks from the new domain
 - The new domain has to protect itself against malicious attacks from the distributed system

Scaling Techniques

- Hiding communication latencies
- Distribution
- Replication

Scaling Techniques

Hiding communication latencies:

- Try to avoid waiting for responses to remote service requests as much as possible
- asynchronous communication (do something else)
- moving part of the computation to the client process

Scaling Techniques

Distribution

- Taking a component, splitting into smaller parts, and spreading these parts across the system
- Example:
 - (1) The World Wide Web
 - (2) Domain Name Service (DNS)
- Hierarchically organized into a tree of domains
- Each domain divided into non overlapping zones
- The names in each domain handled by a single name server

Replication

Caching (client-driven)

- increase availability
- balance the load
- reduce communication latency
- but, consistency problems

Flexibility

Monolithic kernel

The collective kernel structure, Includes file system, directory, and process management

Micro-kernel

- An interprocess communication mechanism
- Some memory management
- A small amount of low-level process management and scheduling
- Low level input/output
- File server
- Directory server
- Process server

Reliability / *Dependability*

- Fault tolerance
 - Redundancy techniques
 - Distributed control
- Fault avoidance
 - Design of components
- Fault detection and recovery
 - Atomic transactions, stateless servers, ack and time out based retransmission.

- Performance
 - Response time, throughput, system utilization, bandwidth requirements
- Global knowledge
 - No shared memory, no global clock, how to find global state